

Parallel Programming

Homework 3: All-Pairs Shortest Path

- 111065524 李懷-

Implementation

a. Which algorithm do you choose in hw3-1?

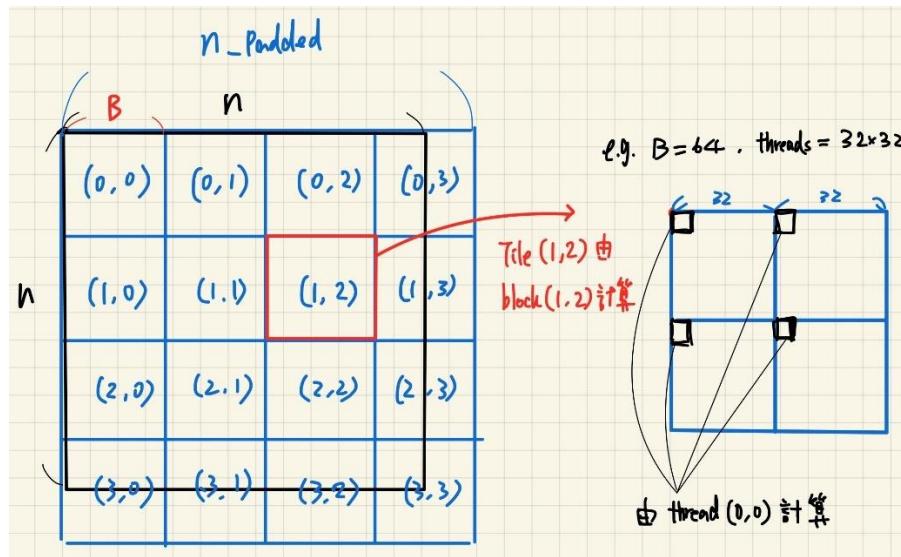
我使用的演算法為 Blocked Floyd-Warshall，針對 CPU 版使用的 threading library 為 openmp，並搭配 vectorization 去處理每個 tile(這裡我用 tile 稱呼避免跟之後 cuda 的 block 搞混)。

b. How do you divide your data in hw3-2, hw3-3?

HW3-2(Single GPU)

將整體 $n * n$ 的 vertices 矩陣切分成 $rounds * rounds$ 個 tiles($rounds = \lceil n/BLOCK_FACTOR \rceil$)，然後將 launch 的 cuda block 一對一 map 到每個 tile，一個 block 設 $THREAD_PER_BLOCK * THREAD_PER_BLOCK$ 個 threads，也就是一個 thread 需負責處理 $(BLOCK_FACTOR / THREAD_PER_BLOCK)^2$ 個 elements，下圖範例為 $BLOCK_FACTOR = 64$ ， $threads = 32 * 32$:

```
#define BLOCK_FACTOR 64 // max shared mem is 49152  
#define THREAD_PER_BLOCK 32
```



下面是將 $n * n$ 的 data 加 padding，這樣在計算每對 vertices 的 distance 時可以不需顧慮當 n 不整除於 block_factor 時邊緣 tile 的 boundary check，即便處理到超出 boundary 的 data，到時 unpad 回來並不會影響結果：

```

int rounds = (n + B - 1) / B;
// pad the matrix size to multiple of B to avoid boundary checks
int n_padded = rounds * B;

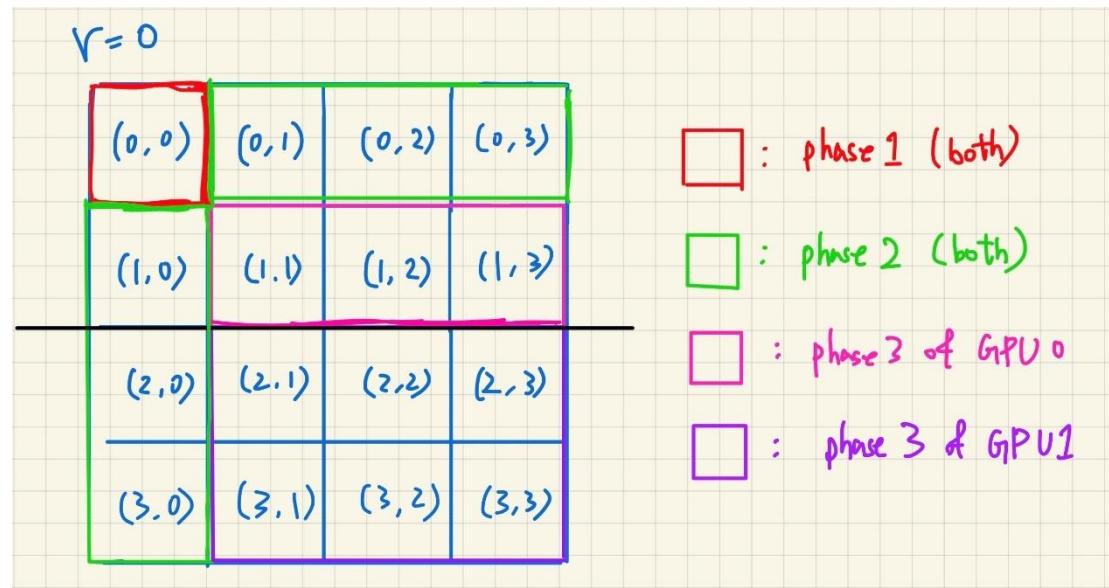
int *h_Dist_padded = (int*)malloc(n_padded * n_padded * sizeof(int));
for (int i = 0; i < n_padded; ++i) {
    for (int j = 0; j < n_padded; ++j) {
        if (i < n && j < n) {
            h_Dist_padded[i * n_padded + j] = (i == j) ? 0 : INF;
        } else {
            // Dist[i][j] = INF;
            h_Dist_padded[i * n_padded + j] = INF;
        }
    }
}

int pair[3];
for (int i = 0; i < m; ++i) {
    fread(pair, sizeof(int), 3, file);
    // Dist[pair[0]][pair[1]] = pair[2];
    h_Dist_padded[pair[0] * n_padded + pair[1]] = pair[2];
}

```

HW3-3(Multi-GPU)

threads 處理 data 的方式跟 single GPU 版的一樣，一樣是一個 block 處理一個 tile，一個 thread 處理 tile 內對應的(*Block_Factor/blockDim.x*)個 datas，tiles 分佈的部分，因為有兩張 GPU，且由於 Blocked Floyd-Warshall 當 vertices 很多時主要的計算都會在 phase3 上，phase1 與 phase2 需要處理的 tiles 相對少，所以與其讓各 GPU 計算自己部份的 phase1 跟 2 後再互傳，增加 communication 成本，直接讓兩張 GPU 都去算當下 round 的 pivot 及 pivot row&column 反而比較快，而在 phase3 時讓 GPU0 負責計算 $n \times n$ 上半，GPU1 負責下半的 tiles，這樣的分法各 GPU 只需在 round 剛開始時同步當下的 row，不需每個 phase 算完就同步一次(communication 的細節放在 1-d)。



c. What's your configuration in hw3-2, hw3-3? And why? (e.g. blocking factor, #blocks, #threads)

HW3-2(Single GPU)

最一开始我 **block_factor** 是设 64，搭配 32×32 的 **#thread**，然后一个 block 一个 tile，所以一个 thread 需处理 4 个 data，但考量到 GTX1080 的 shared memory per block 有 49152 bytes，**block_factor** 设 64 不能完全利用 shared memory，后来我改设 78 ($78 \times 78 \times 2 \times 4 = 48672$ 刚好为最大可设的 **block_factor**)，*2 是因为 phase2 跟 phase3 計算所需的 data dependent 的 tiles 有两塊)，相较于 64 更大的 tile 更能利用的 shared memory(詳細比較放在 optimization)，**#threads** 则改设 39×13 (为了避免 boundary check，所以须可以整除 78)，代表一个 thread 需处理 12 个 data，这降低了 sm warps 的 occupancy，但同时提升了 ilp(Instruction-Level Parallelism) (詳細在 Section 2 的 Profiling Results 有說明):

```
#define BLOCK_FACTOR 78 // max shared mem is 49152
#define THREAD_PER_BLOCK_X 39
#define THREAD_PER_BLOCK_Y 13
```

#blocks 则根据 phase 而有所不同:

Phase1 是处理 $\text{tile}(r, r)$ (r : 第几个 round)，也就是我们只需要 launch 一个 block :

Phase2 是处理以 $\text{tile}(r, r)$ 为中心的十字上的 tiles，所需的 blocks 数为 $2 * (\text{rounds} - 1)$;

Phase3 则是处理以除此之外的 tiles，所需的 blocks 数为 $(\text{rounds} - 1)^2$;

由於我直接让 block index 与 tile index 对应，所以 phase2 跟 phase3 的 kernel 我是直接 launch

rounds^2 个 blocks，并在 kernel 中 return 掉当下 phase 无需计算的 tile 的 index 即可:

```
__global__ void phase2_cal(int *s, int n, int rounds, int r) {
    int b_i = blockIdx.y; // tile row index
    int b_j = blockIdx.x; // tile column index
    bool is_pivot_row = (b_i == r && b_j != r);
    bool is_pivot_col = (b_j == r && b_i != r);
    if (!is_pivot_row && !is_pivot_col)
        return;
}

__global__ void phase3_cal(int *s, int n, int rounds, int r) {
    int b_i = blockIdx.y; // tile row index
    int b_j = blockIdx.x; // tile column index
    if (b_i == r || b_j == r)
        return;
}
```

Kernel launch 的部分:

```
for (int r = 0; r < rounds; ++r) {
    phase1_cal<<<1, threadsPerBlock>>>(d_Dist, n_padded, B, r);
    phase2_cal<<<dim3(rounds, rounds), threadsPerBlock>>>(d_Dist, n_padded, rounds, r);
    phase3_cal<<<dim3(rounds, rounds), threadsPerBlock>>>(d_Dist, n_padded, rounds, r);
}
```

實測後(**block_factor**, **#thread**)設(78, 39×13)能 pass 較多 testcases:

HW3-3(Multi-GPU)

(**block_factor**, #**thread**) 跟 hw3-2 一樣設(78, 39*13) , Kernel 一樣是 launch rounds^2 個 blocks , 由於兩張 GPU 各只處理一半的 data(我是分成上下兩半) , 因此在 phase3 kernel 裡會把在處理範圍外無需計算的 block 擋掉:

```
if (b_i == r || b_j == r)
    return;

if (b_i < start_row || b_i >= end_row )
    return;
```

```
phase1_cal<<<1, threadsPerBlock>>>(d_Dist[tid], n_padded, B, r);

phase2_cal<<<dim3(rounds, rounds), threadsPerBlock>>>(d_Dist[tid], n_padded, rounds, r, tid);

phase3_cal<<<dim3(rounds, rounds), threadsPerBlock>>>(d_Dist[tid], n_padded, rounds, r, tid, start_row ,end_row);
```

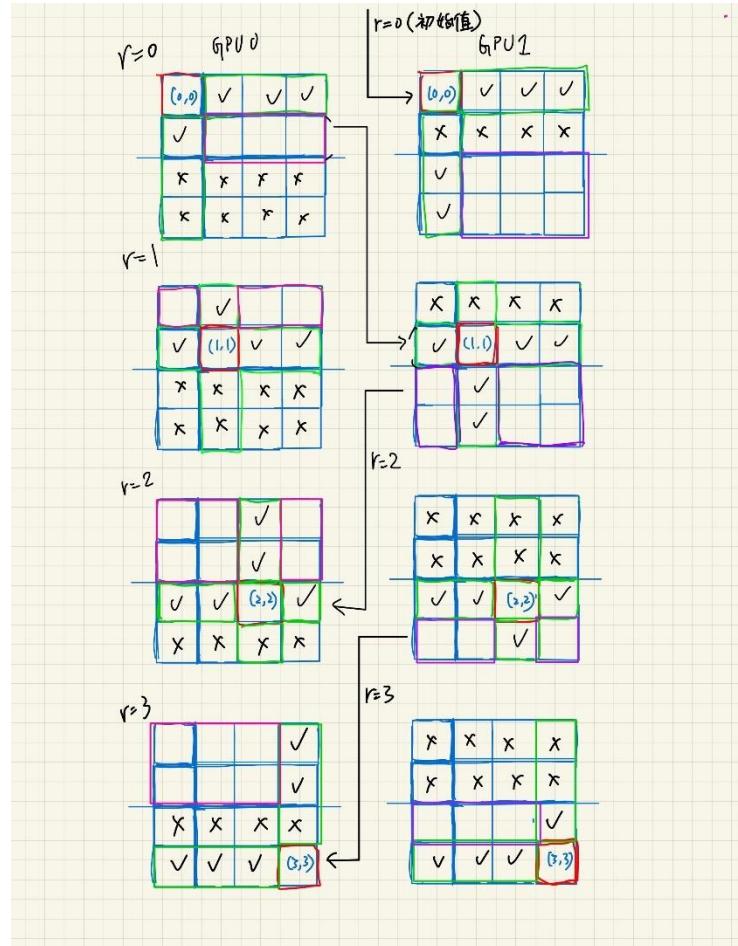
d. How do you implement the communication in hw3-3?

HW3-3(Multi-GPU)

首先是用 openmp create 兩個 threads , 一張 GPU map 到一個 thread , data 交換最一開始我是使用 zero-copy 的方式讓多張 GPU 來讀取同段 pinned 的 host memory 來實作，將 phase2 及 phase3 的 tiles 平分給 GPU , 各 GPU 在 kernel 中計算完會將自己負責的 tiles 直接寫到 pinned 的 host memory , 這樣可以不需要 device 間 cudaMemcpy , 且 workload 可以更平均，但由於讀寫 host memory 太慢了，所以效率差，後來改成用將 $n \times n$ 的 data 水平切分成兩半，phase3 時 GPU0 處理上半，GPU1 處理下半，並用 cudaMemcpyPeer 的方式在每個 round 開始時同步當下的 pivot row 。

各 GPU 在每個 round 計算 phase3 時，只需當下的 pivot row 及其處理部分對應到的 pivot column tiles 即可，例: 在 $r=0$ 時，GPU0 計算上半的 phase3 tiles 需要整條 pivot row，但 pivot column 只需要上半的即可，GPU1 則是只需要整條 pivot row 及下半的 pivot column，且由於每個 round 在計算 pivot 及 pivot row&column 時，需要前一個 round 的 phase3 的結果，所以需在每個 round 的 phase1 前將當下的 row 同步，例如 4×4 個 tiles，共 4 個 rounds，在 $r=0,1$ 時，pivot row 會是在上半，所以 GPU0 會將第 1, 2 列傳給 GPU1 讓 GPU1 可以計算正確的 pivot 及 pivot row&column；當進入到 $r=2,3$ 時，pivot row 會是在下半，所以 GPU1 會將第 3,4 列傳給 GPU0，透過這種方式能確保，每個 round 各 GPU 會用到的 tile 都為最新 update 的，最後 GPU0 將其上半，GPU1 將其下半各別寫回 host memory 即可。

打勾的 tile 代表各 GPU 在 phase3 時 dependent 的 tile，也就是只要確保在每個 round 開始時的 pivot row 都是最新的，就能確保 phase3 是正確的，各 GPU 會用到的 pivot column 的 tiles 由於都位在自己處理的 rows 內，所以一定能確保是最新：



e. Briefly describe your implementations in diagrams, figures or sentences.

HW3-1(CPU)

使用 OpenMP 平行化外層的兩個 loops，也就是分配 tiles 給不同的 threads 去執行，每個 threads 負責處理 tile 內部的運算，並搭配 128bit SIMD register 可一次處理 4 個 INT，在 CPU 版本下，block factor 越小搭配 omp scheduler 有越好的 load balance，因為 threads 處理完一小塊 tile 後便可執行下一個。

```

#pragma omp parallel for num_threads(num_threads) schedule(static) collapse(2)
// #pragma omp parallel for schedule(dynamic, chunk_size)
for (int b_i = block_start_x; b_i < block_end_x; ++b_i) {
    for (int b_j = block_start_y; b_j < block_end_y; ++b_j) {
        // To calculate B*B elements in the block (b_i, b_j)
        // For each block, it need to compute B times, data dependent
        for (int k = Round * B; k < (Round + 1) * B && k < n; ++k) {

for (int i = block_internal_start_x; i < block_internal_end_x; ++i) {
    __m128i dist_ik = _mm_set1_epi32(Dist[i][k]);
    int j;
    for (j = block_internal_start_y; j + 3 < block_internal_end_y; j += 4) {
        __m128i dist_kj = _mm_loadu_si128((__m128i*)&Dist[k][j]);
        __m128i dist_ij = _mm_loadu_si128((__m128i*)&Dist[i][j]);
        __m128i sum = _mm_add_epi32(dist_ik, dist_kj);
        __m128i mask = _mm_cmplt_epi32(sum, dist_ij);
        __m128i result = _mm_or_si128(_mm_and_si128(mask, sum), _mm_andnot_si128(mask, dist_ij));
        _mm_storeu_si128((__m128i*)&Dist[i][j], result);
    }
}

```

HW3-2(Single-GPU)

Shared memory assignment:

分成 phase1,2,3 共 3 個 kernels，每個 phase 的各 block 會先將其需要用到的 tiles 從 global memory load 進其 block 的 shared memory，而每個 block 的 thread 各負責處理 $(BLOCK_FACTOR/THREAD_PER_BLOCK)^2$ 個 datas，對應方式請參考 1-b。

Phase 1:

```
#pragma unroll
for (int i = 0; i < MAX2(BLOCK_FACTOR/THREAD_PER_BLOCK, 1); ++i) [
    #pragma unroll
    for (int j = 0; j < MAX2(BLOCK_FACTOR/THREAD_PER_BLOCK, 1); ++j) {
        S[i*THREAD_PER_BLOCK+ty][j*THREAD_PER_BLOCK+tx] = s[(i * THREAD_PER_BLOCK + gy) * n + (j * THREAD_PER_BLOCK + gx)];
    }
]
__syncthreads();
```

Phase2:

```

#pragma unroll
for (int i = 0; i < MAX2(BLOCK_FACTOR/THREAD_PER_BLOCK, 1); ++i) {
    #pragma unroll
    for (int j = 0; j < MAX2(BLOCK_FACTOR/THREAD_PER_BLOCK, 1); ++j) {
        S[i * THREAD_PER_BLOCK + ty][j * THREAD_PER_BLOCK + tx] = s[(i * THREAD_PER_BLOCK + gy) * n + (j * THREAD_PER_BLOCK + gx)];
        S_pivot[i * THREAD_PER_BLOCK + ty][j * THREAD_PER_BLOCK + tx] = s[(r * B + i * THREAD_PER_BLOCK + ty) * n + (r * B + j * THREAD_PER_BLOCK + gx)];
    }
}
syncthreads();

```

Phase3:

```

int local_vals[MAX2(BLOCK_FACTOR/THREAD_PER_BLOCK, 1)][MAX2(BLOCK_FACTOR/THREAD_PER_BLOCK, 1)];

#pragma unroll
for (int i = 0; i < MAX2(BLOCK_FACTOR/THREAD_PER_BLOCK, 1); i += 1) {
    #pragma unroll
    for (int j = 0; j < MAX2(BLOCK_FACTOR/THREAD_PER_BLOCK, 1); j += 1) {
        // if ((pivot_row_start_x + i) < n && (j+pivot_row_start_y) < n)
        S_pivot_row[i*THREAD_PER_BLOCK+ty][j*THREAD_PER_BLOCK+tx] = s[(pivot_row_start_x + i*THREAD_PER_BLOCK + ty) * n + (j*THREAD_PER_BLOCK+tx)];
        // if ((pivot_col_start_x + i) < n && (j+pivot_col_start_y) < n)
        S_pivot_col[i*THREAD_PER_BLOCK+ty][j*THREAD_PER_BLOCK+tx] = s[(pivot_col_start_x + i*THREAD_PER_BLOCK + ty) * n + (j*THREAD_PER_BLOCK+tx)];
        local_vals[i][j] = s[(i*THREAD_PER_BLOCK + gy) * n + (j*THREAD_PER_BLOCK+ gx)];
    }
}
__syncthreads();

```

各 phase 從 shared memory load 資料運算當下的 tile，最後寫回 global memory， (ty, tx) 為 thread 的 index， (gy, gx) 為各 data 在 $n * n$ matrix 的 global index。

Phase1:

```
for (int k = 0; k < BLOCK_FACTOR; ++k) {
    __syncthreads();
    #pragma unroll
    for (int i = 0; i < MAX2(BLOCK_FACTOR/THREAD_PER_BLOCK_Y, 1); ++i) {
        #pragma unroll
        for (int j = 0; j < MAX2(BLOCK_FACTOR/THREAD_PER_BLOCK_X, 1); ++j) {
            int gi = i * THREAD_PER_BLOCK_Y + ty;
            int gj = j * THREAD_PER_BLOCK_X + tx;
            S[gi][gj] = min(S[gi][k] + S[k][gj], S[gi][gj]);
        }
    }

    #pragma unroll
    for (int i = 0; i < MAX2(BLOCK_FACTOR/THREAD_PER_BLOCK_Y, 1); ++i) {
        #pragma unroll
        for (int j = 0; j < MAX2(BLOCK_FACTOR/THREAD_PER_BLOCK_X, 1); ++j) {
            S[(gy + i * THREAD_PER_BLOCK_Y)*n + (gx + j * THREAD_PER_BLOCK_X)] = S[i*THREAD_PER_BLOCK_Y+ty][j*THREAD_PER_BLOCK_X+tx];
        }
    }
}
```

由於 phase1 是 self-dependent，且 k 的 iterations 是 dependent 的，所以要確保整個 tile 裡的 data 在每個 k 時都是用到最新 update 的值，也因為我們 launch 的 block threads 數大於 32，一個 block 有多個 warps，所以未必免有 warps 先進到下個 iterations 用到尚未 update 的資料，所以在每個 k 都用 `syncthreads()` 確保 block 裡的 threads 在每個 k 時是 sync 的。

Phase2:

```
#pragma unroll
for (int k = 0; k < BLOCK_FACTOR; ++k) {
    #pragma unroll
    for (int i = 0; i < MAX2(BLOCK_FACTOR/THREAD_PER_BLOCK_Y, 1); ++i) {
        #pragma unroll
        for (int j = 0; j < MAX2(BLOCK_FACTOR/THREAD_PER_BLOCK_X, 1); ++j) {
            int gi = i * THREAD_PER_BLOCK_Y + ty;
            int gj = j * THREAD_PER_BLOCK_X + tx;
            if (is_pivot_col) {
                // S[i][j] = min(S[i][j], S[i][k] + S_pivot[k][j]);
                S[gi][gj] = min(S[gi][gj], S[gi][k] + S_pivot[k][gj]);
            }else{
                S[gi][gj] = min(S[gi][gj], S_pivot[gi][k] + S[k][gj]);
            }
        }
    }

    #pragma unroll
    for (int i = 0; i < MAX2(BLOCK_FACTOR/THREAD_PER_BLOCK_Y, 1); ++i) {
        #pragma unroll
        for (int j = 0; j < MAX2(BLOCK_FACTOR/THREAD_PER_BLOCK_X, 1); ++j) {
            S[(i * THREAD_PER_BLOCK_Y + gy) * n + (j * THREAD_PER_BLOCK_X + gx)] = S[i * THREAD_PER_BLOCK_Y + ty][j * THREAD_PER_BLOCK_X + tx];
        }
    }
}
```

Phase2 在計算時，會根據是 pivot row 或是 col 計算方式會有不同，當 data 在 pivot row 時，其會用到對應到的 pivot tile 的 row(`S_pivot[gi][k]`)以及 self tile 的 column；當 data 在 pivot column 時，則是反過來。

Phase3:

```
#pragma unroll
for (int k = 0; k < BLOCK_FACTOR; ++k) {
    #pragma unroll
    for (int i = 0; i < MAX2(BLOCK_FACTOR/THREAD_PER_BLOCK_Y, 1); ++i) {
        #pragma unroll
        for (int j = 0; j < MAX2(BLOCK_FACTOR/THREAD_PER_BLOCK_X, 1); ++j) {
            int gi = i * THREAD_PER_BLOCK_Y + ty;
            int gj = j * THREAD_PER_BLOCK_X + tx;
            int sum = S_pivot_col[gi][k] + S_pivot_row[k][gj];
            local_vals[i][j] = min(local_vals[i][j], sum);
        }
    }
}

#pragma unroll
for (int i = 0; i < MAX2(BLOCK_FACTOR/THREAD_PER_BLOCK_Y, 1); ++i) {
    #pragma unroll
    for (int j = 0; j < MAX2(BLOCK_FACTOR/THREAD_PER_BLOCK_X, 1); ++j) {
        s[(i * THREAD_PER_BLOCK_Y + gy) * n + (j * THREAD_PER_BLOCK_X + gx)] = local_vals[i][j];
    }
}
```

Phase3 在計算時，會使用到對應到的 pivot row 跟 column 的 tile，local_vals 為暫存當下的 distance，最後算完後再 load 回 global memory，這樣可以避免頻繁寫入到 global memory。

HW3-3(Multi-GPU)

Data Exchange between GPUs:

```
for (int r = 0; r < rounds; ++r) {
    // copy specific row to other GPU, r*BLOCK_FACTOR*n_padded: row start index
    cudaMemcpyPeer(d_Dist[!tid]+(r*BLOCK_FACTOR*n_padded), !tid, d_Dist[tid]+(r*BLOCK_FACTOR*n_padded), tid, (r>=start_row && r<end_row)*BLOCK_FACTOR*n_padded*sizeof(int));
    #pragma omp barrier

    phase1_cal<<<1, threadsPerBlock>>>(d_Dist[tid], n_padded, B, r);
    phase2_cal<<<dim3(rounds, rounds), threadsPerBlock>>>(d_Dist[tid], n_padded, rounds, r, tid, start_row ,end_row);
    phase3_cal<<<dim3(rounds, rounds), threadsPerBlock>>>(d_Dist[tid], n_padded, rounds, r, tid, start_row ,end_row);
}
```

在每個 round 開始時由負責處理當 round 的 pivot row 的 GPU 將 row 用 cudaMemcpyPeer 傳給另一張 GPU。

Phase3 各 GPU 只會處理自己負責的 rows:

```
__global__ void phase3_cal(int *s, int n, int rounds, int r, int tid, int start_row, int end_row) {

    int b_i = blockIdx.y;    // tile row index
    int b_j = blockIdx.x;    // tile column index

    if (b_i == r || b_j == r)
        return;

    if (b_i < start_row || b_i >= end_row )
        return;
```

Profiling Results (hw3-2)

Invocations	Metric Name	Metric Description	Min	Max	Avg
Device "NVIDIA GeForce GTX 1080 (0)"					
Kernel: phase3_cal(int*, int, int, int)					
65	branch_efficiency	Branch Efficiency	100.00%	100.00%	100.00%
65	gld_throughput	Global Load Throughput	264.02GB/s	269.01GB/s	266.91GB/s
65	gst_throughput	Global Store Throughput	111.39GB/s	113.50GB/s	112.61GB/s
65	shared_load_throughput	Shared Memory Load Throughput	3176.6GB/s	3236.8GB/s	3211.5GB/s
65	shared_store_throughput	Shared Memory Store Throughput	224.33GB/s	228.57GB/s	226.79GB/s
65	shared_efficiency	Shared Memory Efficiency	32.01%	32.01%	32.01%
65	inst_integer	Integer Instructions	3113062134	3113062134	3113062134
65	sm_efficiency	Multiprocessor Activity	99.52%	99.67%	99.60%
65	achieved_occupancy	Achieved Occupancy	0.473876	0.476032	0.475110

Fig. 2-1 Block_factor=78, threads=39*13, testcase:c20.1

這裡我用 nvprof 去 profile phase3_cal 這個最大的 kernel， testcase 選擇 c20.1 主要是 profiling 很花時間，太大的 testcase 會跑太久被 slurm terminate 掉。可看到當一個 block 的 thread 設 39*13 時，occupancy 會下降，這是因為一個 block 需要 16 個 warps($\text{ceil}(39*13/32)=16$)，但又因為我們一個 block 用滿了 GTX 1080 的 per block shared memory limit(48KB)，然後 max shared memory per SM 是 96KB，代表在 block_factor =78，threads 數為 39*13 的情況下，一個 SM 只能有兩個 blocks， $2*16/64=0.5$ ，導致較低的 warps occupancy，當我將 threads 數為 32*32 時，兩個 blocks 共 64 warps，所以可以佔滿 occupancy:

Invocations	Metric Name	Metric Description	Min	Max	Avg
Device "NVIDIA GeForce GTX 1080 (0)"					
Kernel: phase3_cal(int*, int, int, int)					
79	branch_efficiency	Branch Efficiency	100.00%	100.00%	100.00%
79	gld_throughput	Global Load Throughput	196.14GB/s	200.87GB/s	198.82GB/s
79	gst_throughput	Global Store Throughput	65.380GB/s	66.957GB/s	66.274GB/s
79	shared_load_throughput	Shared Memory Load Throughput	3138.26GB/s	3213.9GB/s	3181.1GB/s
79	shared_store_throughput	Shared Memory Store Throughput	130.76GB/s	133.91GB/s	132.55GB/s
79	shared_efficiency	Shared Memory Efficiency	70.00%	70.00%	70.00%
79	inst_integer	Integer Instructions	2716608512	2716608512	2716608512
79	sm_efficiency	Multiprocessor Activity	99.50%	99.66%	99.58%
79	achieved_occupancy	Achieved Occupancy	0.931585	0.939131	0.934986

Fig. 2-2 Block_factor=64, threads=32*32, testcase:c20.1

但較低的 occupancy 不一定代表較差的 optimization，Nvidia 官方 CUDA document 中有提到「Higher occupancy does not always equate to higher performance-there is a point above which additional occupancy does not improve performance.」，occupancy 只要足夠可以掩蓋掉 memory stall 的 latency 即可，而在使用 shared memory latency 很低的情況下，過高的 occupancy 反而沒有必要，因此我們可以降低 threads 數犧牲一點 occupancy，讓每個 thread 處理更多 data 提高 ILP(Instruction-Level Parallelism)，這樣可以更好地利用每個 thread。另外，Fig. 2-1 的 shared store/global load/store 皆明顯上升，主要也是因為一個 thread 一次處理處理更多資料，但可以看到 shared load 跟 shared efficiency 有下降，主要是因為 bank conflicts。

Experiment & Analysis (Nvidia GPU)

a. System Spec:

使用的機器為這門課提供的 Apollo Server。

b. Blocking Factor (hw3-2)

由於 GTX 1080 的 shared memory per block 的上限，針對 shared memory 的實驗上限只能設 78，總共試了以下幾種 BF:{16, 32, 48, 64, 78}，針對 global memory 的實驗則是 BF={16, 32, 64, 128, 256}， testcase 為 c20.1:

```
for BF in "${BLOCK_FACTORS[@]}"; do
    echo "=====
    echo "BLOCK_FACTOR = $BF"
    echo ====="

    make -B "$EXPERIMENT" BLOCK_FACTOR=$BF

    # Run nvprof, but keep *output* file distinct
    srun -p nvidia -N1 -n1 --gres=gpu:1 nvprof --csv --log-file ${OUTDIR}/raw.csv \
        | -f -o ${OUTDIR}/time_${BF}.nvprof --print-summary --csv --normalized-time-unit ms \
        --kernels phase3_cal \
        ./${EXPERIMENT} ../../testcases/${TEST_CASE} ${TEST_CASE}.out

    grep '^"' ${OUTDIR}/raw.csv > ${OUTDIR}/time_${BF}.csv && rm ${OUTDIR}/raw.csv

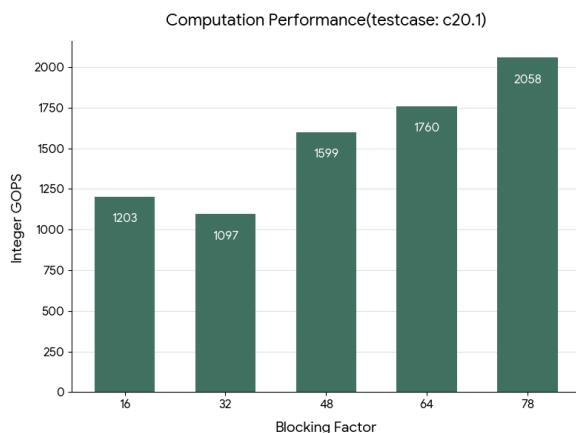
    srun -p nvidia -N1 -n1 --gres=gpu:1 nvprof --csv --log-file ${OUTDIR}/raw.csv \
        -f -o ${OUTDIR}/metrics_${BF}.nvprof --print-summary --csv \
        --kernels phase3_cal \
        -m $NVPROF_METRIC \
        ./${EXPERIMENT} ../../testcases/${TEST_CASE} ${TEST_CASE}.out

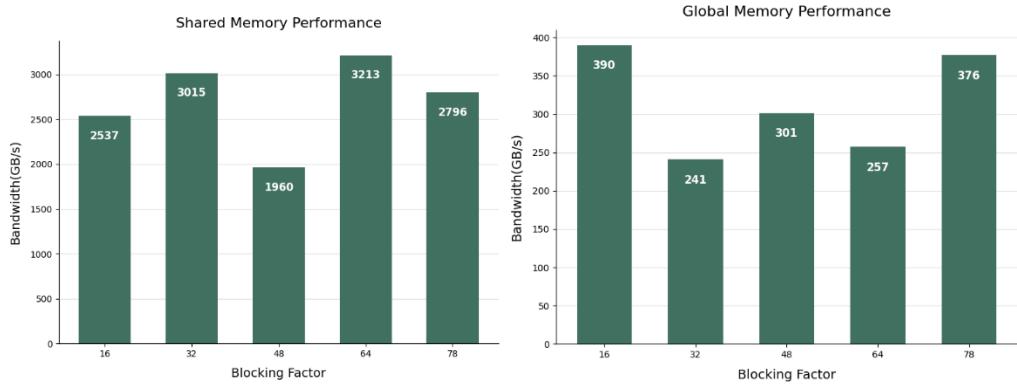
    grep '^"' ${OUTDIR}/raw.csv > ${OUTDIR}/metrics_${BF}.csv && rm ${OUTDIR}/raw.csv
done
```

Integer GOPS 的算法我是用 nvprof 去紀錄各 kernel 一個 invocation 平均的時間 t_{avg} (ms)，再用 inst_integer metric 得出的一個 invocation 平均的 int instructions I_{avg} ，

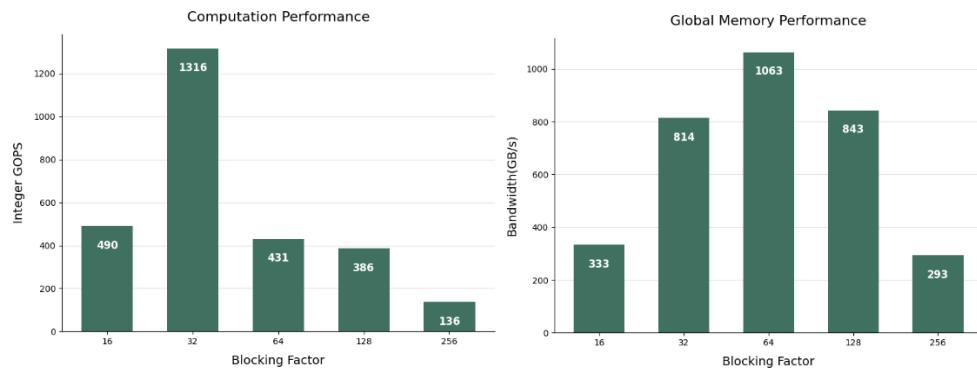
用以下算式得出: $Int\ GOPS = \frac{I_{avg}}{(t_{avg}(ms) * 10^{-3})} * 10^{-9}$

Using shared mem:





W/O shared mem(coalesced global memory access):



可以看到有用 shared memory 的情況下，int GOPS 基本上是隨著 BF 上升的，但因為 shared memory per block 的上限，最高只能設 78，再來針對 BF 78，因為 threads 數沒辦法設 39×39 (維持一個 thread 處理 4 個 datas)，所以改設 39×13 ，occupancy 下降，所以 shared memory 得 bandwidth 有下降，但 ILP 上升，所以仍可以提升整體的 computation performance；沒有使用 shared memory 的情況下可以將 BF 設更高，但過大的 BF 會導致整體效率下降，主要是因為高 ILP 在 global memory 的 bottleneck 下，使 stall latency 更高，即便高 occupancy(threads 數皆為 32×32)也無法掩蓋掉這些 latency。

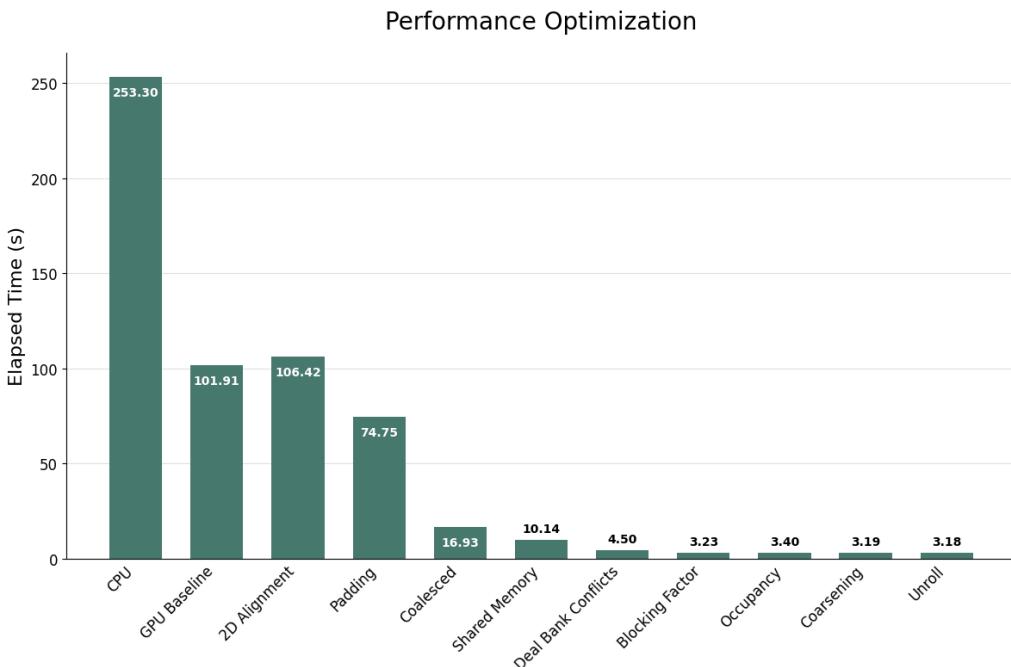
c. Optimization(hw3-2):

由於太大的 testcase 在跑 baseline 時會超過 slurm 每個 task 5 分鐘限制，所以我選 p12k1 作為實驗的 testcase。

我依序試了以下幾種 optimizations，optimizations 是 cumulative 的，所以每一步都是在前一項優化的基礎上繼續優化：

1. 2D alignment: grid 跟 block 由原本 baseline 的 1D 改成 2D；
2. Padding: 原本 $n \times n$ 的 data 改成 $n_{\text{padded}} \times n_{\text{padded}}$ 可以整除 BLOCK_FACTOR，移除 index 的 boundary check；
3. Coalesced memory access: 將讀取 global memory 的 index 由 col-major 改為 row-major，確保一個 warp 是讀取連續區塊的 memory；

4. Shared memory(w/ bank conflicts): 因為演算法會一直 reuse dependent tile 的 data，所以將每個 phase 會用到的 tile load 到 shared memory 上再進行計算，但用 col major 去 load/store shared memory，所以有嚴重 bank conflict;
5. Shared memory(deal with bank conflicts): 改用 row-major 去 load/store shared memory，降低 bank conflict;
6. Blocking factor tuning: 將 BF 由 32 提高至 64；
7. Occupancy optimization: threads 數由 $16*16$ 提高至 $32*32$ ，由於提高一個 block 的 shared memory usage 後，sm 上最多能同時有的 blocks 數下降，此時若 threads 數仍為 $16*16$ 的話，occupancy 低，提升 threads 數可提高 sm 上的 warps 數以提高 occupancy；
8. Thread coarsening: 將一個 block 的 threads 數由 $32*32$ 降為 $32*8$ ，此方法降低了 sm 上 warps 的 occupancy，但讓一個 thread 一次計算更多 data，可提高 ILP;
9. Loop unroll: 將 for (int k = 0; k < BLOCK_FACTOR; ++k) 此 for loop 用 openacc directive 去 unroll，減少 iterations 數；



接著我用 nvprof events 去查看在使用 shared memory 下是否有 bank conflict 發生:

w/ bank conflicts(BF=32, threads=16*16):

Invocations	Event Name	Min	Max	Avg	Total
===== Profiling result:					
===== Event result:					
Device "NVIDIA GeForce GTX 1080 (0)"					
Kernel: phase3_cal(int*, int, int, int)					
157	shared_ld_bank_conflict	87220224	87220224	87220224	$1.3694e+10$
157	shared_st_bank_conflict	23362560	23362560	23362560	3667921920
Kernel: phase1_cal(int*, int, int, int)					
157	shared_ld_bank_conflict	30720	30720	30720	4823040
157	shared_st_bank_conflict	15840	15840	15840	2486880
Kernel: phase2_cal(int*, int, int, int)					
157	shared_ld_bank_conflict	3744000	3744000	3744000	587808000
157	shared_st_bank_conflict	5091840	5091840	5091840	799418880

Invocations	Metric Name	Metric Description	Min	Max	Avg
Device "NVIDIA GeForce GTX 1080 (0)"					
Kernel: phase3_cal(int*, int, int, int)					
157	gld_throughput	Global Load Throughput	64.860GB/s	76.348GB/s	75.502GB/s
157	gst_throughput	Global Store Throughput	21.620GB/s	25.449GB/s	25.167GB/s
157	shared_load_throughput	Shared Memory Load Throughput	3113.3GB/s	3664.7GB/s	3624.1GB/s
157	shared_store_throughput	Shared Memory Store Throughput	691.84GB/s	814.37GB/s	805.36GB/s
157	shared_efficiency	Shared Memory Efficiency	6.25%	6.25%	6.25%
157	achieved_occupancy	Achieved Occupancy	0.951916	0.956411	0.954746

deal with bank conflicts(BF=32, threads=16*16):

pp25s023@apollo-login:~/hw3/experiments/optimization\$ nvprof -i nvprof/hw3-2_no_bank_c20.1.nvprof					
===== Profiling result:					
===== Event result:					
Invocations					
Device "NVIDIA GeForce GTX 1080 (0)"	Event Name	Min	Max	Avg	Total
Kernel: phase3_cal(int*, int, int, int)					
157	shared_ld_bank_conflict	0	0	0	0
157	shared_st_bank_conflict	1557504	1557504	1557504	244528128
Kernel: phase1_cal(int*, int, int, int)					
157	shared_ld_bank_conflict	2048	2048	2048	321536
157	shared_st_bank_conflict	1056	1056	1056	165792
Kernel: phase2_cal(int*, int, int, int)					
157	shared_ld_bank_conflict	249600	249600	249600	39187200
157	shared_st_bank_conflict	339456	339456	339456	53294592
===== Metric result:					
Invocations	Metric Name	Metric Description	Min	Max	Avg
Device "NVIDIA GeForce GTX 1080 (0)"					
Kernel: phase3_cal(int*, int, int, int)					
157	gld_throughput	Global Load Throughput	271.88GB/s	288.37GB/s	276.08GB/s
157	gst_throughput	Global Store Throughput	90.627GB/s	93.457GB/s	92.026GB/s
157	shared_load_throughput	Shared Memory Load Throughput	2175.1GB/s	2243.0GB/s	2208.6GB/s
157	shared_store_throughput	Shared Memory Store Throughput	362.51GB/s	373.83GB/s	368.10GB/s
157	shared_efficiency	Shared Memory Efficiency	39.29%	39.29%	39.29%
157	achieved_occupancy	Achieved Occupancy	0.877593	0.881395	0.879688
Kernel: phase1_cal(int*, int, int, int)					

可以看到用 row-major 去 access shared memory 可以大幅降低 bank conflicts(由於 threads 是 16*16 所以仍有 2-way bank conflict)，shared memory efficiency 顯著提升。

No bank conflicts(BF=32, threads=32*32):

pp25s023@apollo-login:~/hw3/experiments/optimization\$ nvprof -i nvprof/hw3-2_no_bank_c20.1.nvprof					
===== Profiling result:					
===== Event result:					
Invocations					
Device "NVIDIA GeForce GTX 1080 (0)"	Event Name	Min	Max	Avg	Total
Kernel: phase3_cal(int*, int, int, int)					
157	shared_ld_bank_conflict	0	0	0	0
157	shared_st_bank_conflict	0	0	0	0
Kernel: phase1_cal(int*, int, int, int)					
157	shared_ld_bank_conflict	0	0	0	0
157	shared_st_bank_conflict	0	0	0	0
Kernel: phase2_cal(int*, int, int, int)					
157	shared_ld_bank_conflict	0	0	0	0
157	shared_st_bank_conflict	0	0	0	0
===== Metric result:					
Invocations	Metric Name	Metric Description	Min	Max	Avg
Device "NVIDIA GeForce GTX 1080 (0)"					
Kernel: phase3_cal(int*, int, int, int)					
157	gld_throughput	Global Load Throughput	181.19GB/s	185.51GB/s	182.71GB/s
157	gst_throughput	Global Store Throughput	60.395GB/s	61.836GB/s	60.905GB/s
157	shared_load_throughput	Shared Memory Load Throughput	2899.0GB/s	2968.1GB/s	2923.4GB/s
157	shared_store_throughput	Shared Memory Store Throughput	128.79GB/s	123.67GB/s	121.81GB/s
157	shared_efficiency	Shared Memory Efficiency	70.00%	70.00%	70.00%
157	achieved_occupancy	Achieved Occupancy	0.915933	0.917798	0.917161
Kernel: phase1_cal(int*, int, int, int)					

使 blockDim.x=32 可以完全移除 bank conflicts。

另外在做 Thread coarsening 時，例如將 block 由 32*32 降為 16*16 threads，以目前讀取 shared memory 的方式，一個 warp 會跨 2 個 rows，此時 warp 中可能會有些 threads 讀到同一個 bank，造成 2-way bank conflict，所以我將 16*16 調整為 32*8，threads 數一

樣，一個 thread 都是處理 16 個 datas，ILP 一樣，但一個 warp 就是一個 row，不會有 bank conflict：

BF=64, Threads=16*16:

Profiling result:						
Event result:		Event Name	Min	Max	Avg	Total
Invocations						
Device "NVIDIA GeForce GTX 1080 (0)"						
Kernel: phase3_cal(int*, int, int, int)						
79	shared_ld_bank_conflict	0	0	0	0	0
79	shared_st_bank_conflict	1557504	1557504	1557504	1557504	123042816
Kernel: phase1_cal(int*, int, int, int)						
79	shared_ld_bank_conflict	16384	16384	16384	16384	1294336
79	shared_st_bank_conflict	8320	8320	8320	8320	657280
Kernel: phase2_cal(int*, int, int, int)						
79	shared_ld_bank_conflict	818688	818688	818688	818688	64676352
79	shared_st_bank_conflict	1317888	1317888	1317888	1317888	104113152
Device "NVIDIA GeForce GTX 1080 (0)"						
Kernel: phase3_cal(int*, int, int, int)						
79	gld_throughput	Global Load Throughput	219.97GB/s	221.42GB/s	220.78GB/s	
79	gst_throughput	Global Store Throughput	73.323GB/s	73.805GB/s	73.595GB/s	
79	shared_load_throughput	Shared Memory Load Throughput	1759.76GB/s	1771.3GB/s	1766.3GB/s	
79	shared_store_throughput	Shared Memory Store Throughput	293.296GB/s	295.22GB/s	294.38GB/s	
79	shared_efficiency	Shared Memory Efficiency	39.29%	39.29%	39.29%	
79	achieved_occupancy	Achieved Occupancy	0.343040	0.345395	0.343951	
Kernel: phase1_cal(int*, int, int, int)						
79	gld_throughput	Global Load Throughput	330.29MB/s	378.60MB/s	372.63MB/s	
79	gst_throughput	Global Store Throughput	330.29MB/s	378.60MB/s	372.63MB/s	
79	shared_load_throughput	Shared Memory Load Throughput	103.226GB/s	118.31GB/s	116.456GB/s	
79	shared_store_throughput	Shared Memory Store Throughput	41.931GB/s	48.065GB/s	47.386GB/s	
79	shared_efficiency	Shared Memory Efficiency	36.67%	36.67%	36.67%	
79	achieved_occupancy	Achieved Occupancy	0.124900	0.124921	0.124918	
Kernel: phase2_cal(int*, int, int, int)						
79	gld_throughput	Global Load Throughput	24.857GB/s	29.650GB/s	28.411GB/s	
79	gst_throughput	Global Store Throughput	12.429GB/s	14.825GB/s	14.206GB/s	
79	shared_load_throughput	Shared Memory Load Throughput	1516.3GB/s	1808.6GB/s	1733.1GB/s	
79	shared_store_throughput	Shared Memory Store Throughput	1640.66GB/s	1956.9GB/s	1875.26GB/s	
79	shared_efficiency	Shared Memory Efficiency	35.24%	35.24%	35.24%	
79	achieved_occupancy	Achieved Occupancy	0.321680	0.332944	0.327422	

BF=64, Threads=32*8:

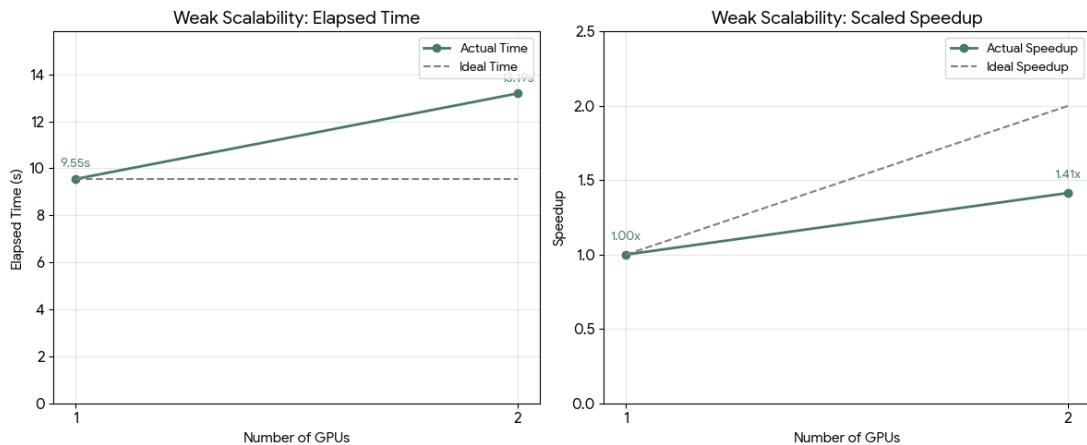
Profiling result:						
Event result:		Metric Name	Metric Description	Min	Max	Avg
Invocations						
Device "NVIDIA GeForce GTX 1080 (0)"						
Kernel: phase3_cal(int*, int, int, int)						
79	shared_ld_bank_conflict	0	0	0	0	0
79	shared_st_bank_conflict	0	0	0	0	0
Kernel: phase1_cal(int*, int, int, int)						
79	shared_ld_bank_conflict	0	0	0	0	0
79	shared_st_bank_conflict	0	0	0	0	0
Kernel: phase2_cal(int*, int, int, int)						
79	shared_ld_bank_conflict	0	0	0	0	0
79	shared_st_bank_conflict	0	0	0	0	0
Invocations		Metric Name	Metric Description	Min	Max	Avg
Device "NVIDIA GeForce GTX 1080 (0)"						
Kernel: phase3_cal(int*, int, int, int)						
79	gld_throughput	Global Load Throughput	220.15GB/s	221.25GB/s	220.64GB/s	
79	gst_throughput	Global Store Throughput	73.382GB/s	73.751GB/s	73.545GB/s	
79	shared_load_throughput	Shared Memory Load Throughput	1761.26B/s	1770.06B/s	1765.16B/s	
79	shared_store_throughput	Shared Memory Store Throughput	146.766GB/s	147.508GB/s	147.09GB/s	
79	shared_efficiency	Shared Memory Efficiency	42.31%	42.31%	42.31%	
79	achieved_occupancy	Achieved Occupancy	0.346492	0.350602	0.348002	
Kernel: phase1_cal(int*, int, int, int)						
79	gld_throughput	Global Load Throughput	329.62MB/s	387.02MB/s	381.66MB/s	
79	gst_throughput	Global Store Throughput	329.62MB/s	387.02MB/s	381.66MB/s	
79	shared_load_throughput	Shared Memory Load Throughput	61.884GB/s	72.566GB/s	71.560GB/s	
79	shared_store_throughput	Shared Memory Store Throughput	20.923GB/s	24.566GB/s	24.226GB/s	
79	shared_efficiency	Shared Memory Efficiency	75.88%	75.88%	75.88%	
79	achieved_occupancy	Achieved Occupancy	0.124930	0.124934	0.124932	
Kernel: phase2_cal(int*, int, int, int)						
79	gld_throughput	Global Load Throughput	32.810GB/s	41.078GB/s	39.648GB/s	
79	gst_throughput	Global Store Throughput	16.405GB/s	20.539GB/s	19.824GB/s	
79	shared_load_throughput	Shared Memory Load Throughput	1394.4GB/s	1745.8GB/s	1685.0GB/s	
79	shared_store_throughput	Shared Memory Store Throughput	1082.76B/s	1355.6GB/s	1308.4GB/s	
79	shared_efficiency	Shared Memory Efficiency	69.21%	69.21%	69.21%	
79	achieved_occupancy	Achieved Occupancy	0.306287	0.320653	0.315491	

可以看到沒了 bank conflicts，但由於對於效能影響最大的是在 phase3 kernel 的 shared load 原本就因 access 的方式沒有造成 bank conflicts，所以即便其他 phase 沒了 bank conflicts，整體效率提升不大，shared memory efficiency 也只有些微上升。

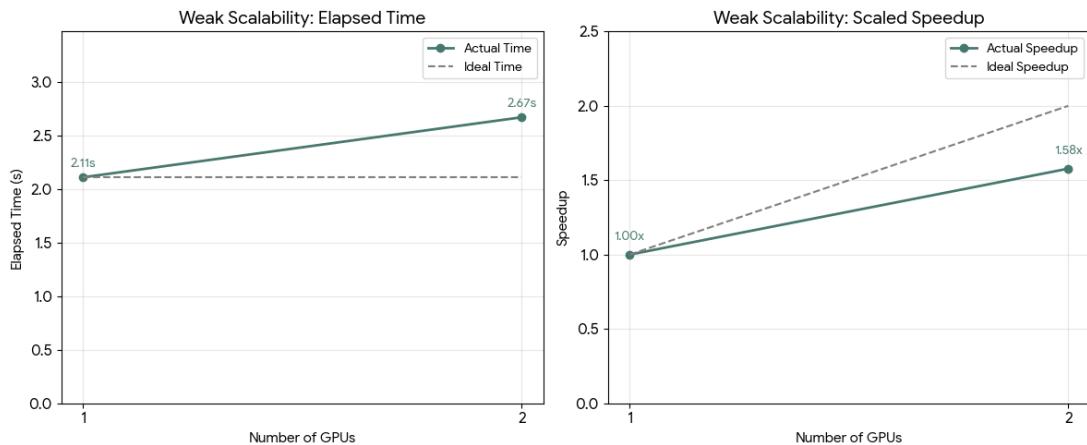
d. Weak scalability (hw3-3)

針對 weak scalability 我採用幾對 testcases，由於 Floyd Warshall 演算法 time complexity 為 $O(N^3)$ ，所以 2 張 GPU 的 testcase 的 vertices 應該要為 1.26 倍，例如 p25k1 的 complexity 約為 p20k1 的 1.95 倍(p20k1 有 20000 個 vertices，p25k1 則有 25000 個 vertices)，可以比較一張及兩張 GPU 下的 weak scalability 及 speedup。

p20k1 v.s. p25k1:



p11k1 v.s. p14k1:



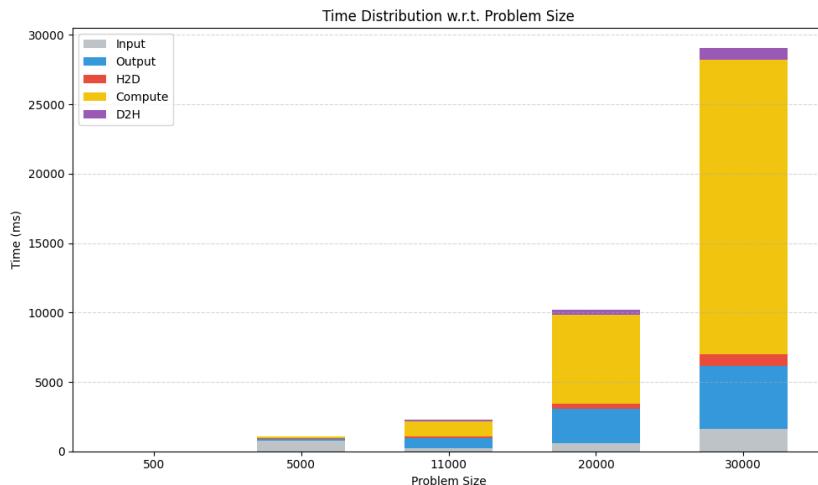
可以看到用 multi-gpu 時因為 GPU 間 communication 的成本很難達到 ideal speedup，且隨著 problem size 越大，speedup 反而變差，主要是因為每次 cudaMemcpy overhead 隨著資料及 rounds 數變大而提高，再來，透過 nvidia-smi 可以看到，兩張 GPUs 是透過 PCIe host bridge 而非像是 PXB, PIX 甚至是更快的 NVlink 直連，而是要經過 CPU，所以即便使用了 cudaMemcpyPeer()，communication overhead 仍很大：

	GPU0	GPU1	CPU Affinity	NUMA Affinity	GPU NUMA ID
GPU0	X	PHB	0-3	0	N/A
GPU1	PHB	X	0-3	0	N/A

e. Time Distribution (hw3-2)

這邊一樣用 nvprof 去 profile 每個 GPU activities 的時間組成，input 跟 output 則另外用 clock_gettime()去紀錄：

```
Type", "Time(%)", "Time", "Calls", "Avg", "Min", "Max", "Name"
"GPU activities", 87.714174, 1006.870458, 142.7.089637, 6.987341, 7.302225, "phase2_cal(int*, int, int, int)"
"GPU activities", 90.924459, 90.964813, 142.0.640597, 0.688454, 0.714696, "phase2_cal(int*, int, int, int)"
"GPU activities", 3.323252, 38.147592, 1, 38.147592, 38.147592, 38.147592, "[CUDA memcpy DtoH]"
"GPU activities", 0.780526, 8.959656, 142.0.063096, 0.061185, 0.069121, "phasel_cal(int*, int, int, int)"
"GPU activities", 0.168075, 1.929334, 1, 1.929334, 1.929334, 1.929334, "init_dist(int*, int)"
"GPU activities", 0.052485, 0.602471, 1, 0.602471, 0.602471, 0.602471, "[CUDA memcpy HtoD]"
"GPU activities", 0.037029, 0.425060, 1, 0.425060, 0.425060, 0.425060, "dist_fill(int*, int, int, int)"
"API calls", 75.597361, 1143.268602, 1, 1143.268602, 1143.268602, 1143.268602, "cudaMemcpy2D"
"API calls", 19.826541, 299.839322, 2, 149.919661, 78.772823, 221.111089, "cudaHostAlloc"
"API calls", 4.178000, 63.184429, 2, 31.592214, 0.924849, 62.259580, "cudaFreeHost"
"API calls", 0.164313, 2.484931, 1, 2.484931, 2.484931, 2.484931, "cudaFree"
"API calls", 0.151476, 2.290792, 428.0.005352, 0.003463, 0.524879, "cudaLaunchKernel"
"API calls", 0.043479, 0.657543, 1, 0.657543, 0.657543, 0.657543, "cudaMemcpy"
"API calls", 0.026415, 0.399471, 2, 0.199735, 0.076017, 0.323454, "cudaMalloc"
"API calls", 0.010532, 0.159276, 114.0.001397, 0.000099, 0.067807, "cuDeviceGetAttribute"
"API calls", 0.001036, 0.015669, 1, 0.015669, 0.015669, 0.015669, "cuDeviceGetName"
"API calls", 0.000610, 0.009218, 1, 0.009218, 0.009218, 0.009218, "cuDeviceGetPCIBusId"
"API calls", 0.000094, 0.001420, 3, 0.000473, 0.000108, 0.001175, "cuDeviceGetCount"
"API calls", 0.000059, 0.000899, 2, 0.000449, 0.000154, 0.000745, "cuDeviceGet"
"API calls", 0.000042, 0.000632, 1, 0.000632, 0.000632, 0.000632, "cuDeviceTotalMem"
"API calls", 0.000027, 0.000401, 1, 0.000401, 0.000401, 0.000401, "cuModuleGetLoadingMode"
"API calls", 0.000016, 0.000241, 1, 0.000241, 0.000241, 0.000241, "cuDeviceGetUuid"
```



可以看到當 problem size 很小時，input 的時間佔比大，主要是因為 GPU kernel 很快就計算完了，隨著 problem size 越大，主要的 bottleneck 來到了 computation。

f. Others:

主要是針對 IO, H2D 及 D2H 做一些細微的優化：

1. 用 cudaMemcpy2D 去 unpad 結果；
2. 用 cudaMallocHost 將 host memory pinned 住來達到較好的 H2D/D2H 的速度；
3. 改用 GPU kernel 去平行化 n*n data matrix 的 initialize 跟 assignment

```

__global__ void init_dist(int *s, int n) {
    int b_i = blockIdx.y; // tile row index
    int b_j = blockIdx.x; // tile column index

    int ty = threadIdx.y;
    int tx = threadIdx.x;

    int tid = ty * blockDim.x + tx;

    if (tx >= BLOCK_FACTOR || ty >= BLOCK_FACTOR) return;

    int block_internal_start_y = b_i * BLOCK_FACTOR;
    int block_internal_start_x = b_j * BLOCK_FACTOR;

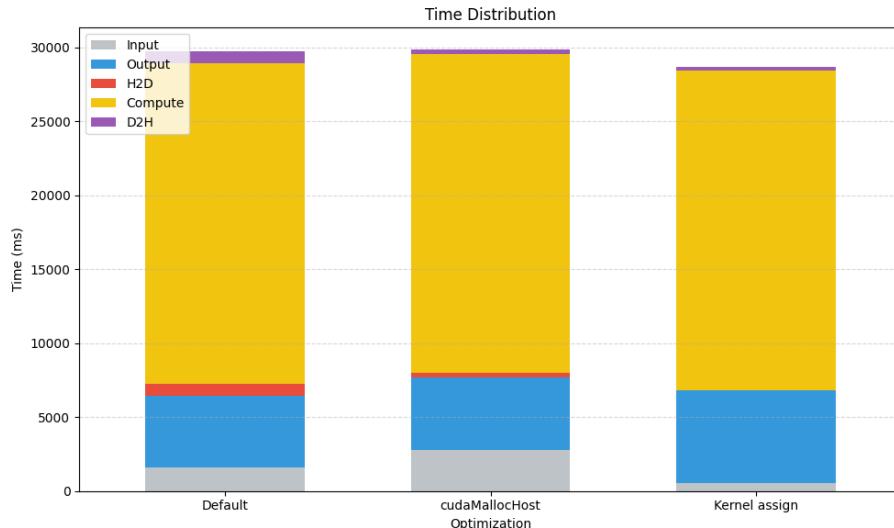
    int gx = block_internal_start_x + tx;
    int gy = block_internal_start_y + ty;

    #pragma unroll
    for (int i = 0; i < MAX2(BLOCK_FACTOR/THREAD_PER_BLOCK_Y, 1); ++i) {
        #pragma unroll
        for (int j = 0; j < MAX2(BLOCK_FACTOR/THREAD_PER_BLOCK_X, 1); ++j) {
            int gi = i * THREAD_PER_BLOCK_Y + gy;
            int gj = j * THREAD_PER_BLOCK_X + gx;
            s[gi * n + gj] = (gi == gj) ? 0 : INF;
        }
    }
}

__global__ void dist_fill(int* s, int *pairs, int n, int m) {
    int idx = blockIdx.x * blockDim.x + threadIdx.x;
    if (idx < m) {
        int u = pairs[3 * idx];
        int v = pairs[3 * idx + 1];
        int w = pairs[3 * idx + 2];
        s[u * n + v] = w;
    }
}

```

這裡透過大的 testcase(p30k1) 較能看出 3 所帶來的 input time 的差異：



AMD GPU Porting and Analysis

a. Implementation

Explain how you ported the CUDA code to HIP code. (hw3-2):

基本上 implementation 跟 hw3-2 cuda 版本的一樣，只是將 cuda 的 function 替換成 hip，例
cudaMalloc -> hipMalloc。

b. Experiment & Analyze

我是用 rocprof 去做 profiling，metrics 有 SQ_INSTS_VMEM_WR

SQ_INSTS_VMEM_RD SQ_INSTS_VALU SQ_LDS BANK_CONFLICT SQ_INSTS_LDS，由於 rocprof 沒有直接提供 shared/global memory throughput，所以只能透過 SQ_INSTS_LDS 跟

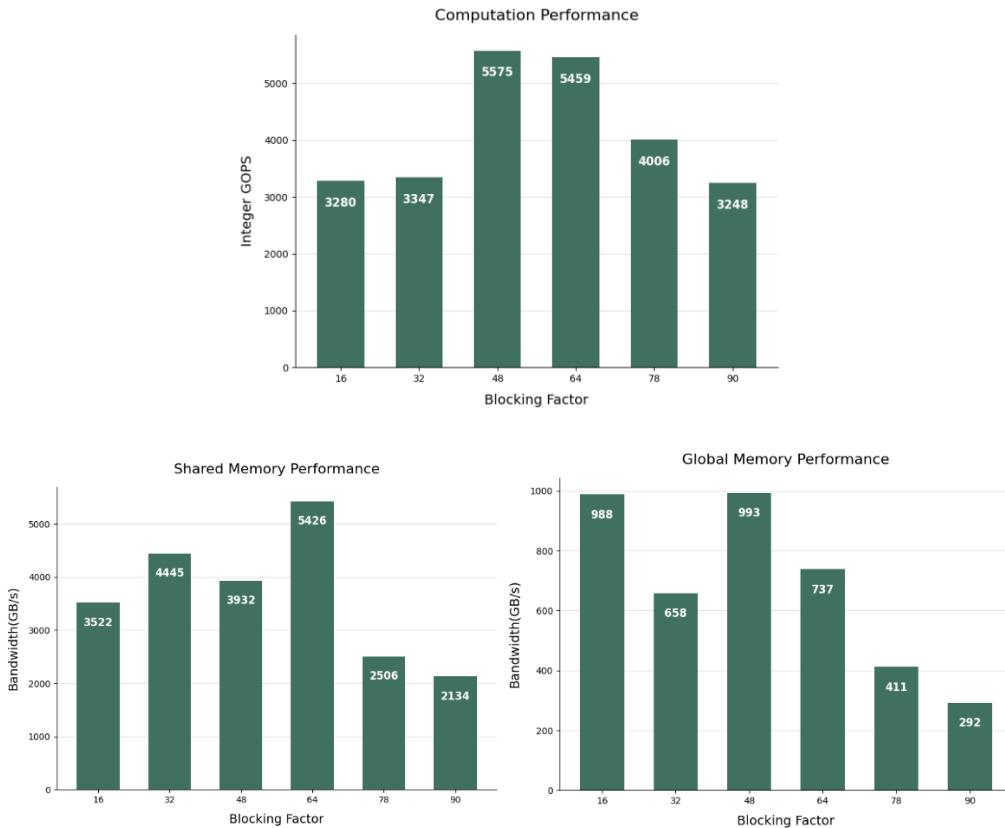
SQ_INSTS_VMEM_RD/WR 及 durations 去算，各算法如下：

$$Int\ GOPS = \frac{\sum_{i=1}^3 I_{phase_i_avg}}{\sum_{i=1}^3 t_{phase_i_avg} (ns) * 10^{-9}} * 10^{-9}$$

$$Bandwidth = \frac{\sum_{i=1}^3 R_{phase_i_avg}}{\sum_{i=1}^3 t_{phase_i_avg} (ns) * 10^{-9}} * 10^{-9}$$

1. $I_{phase_i_avg}$: 各 phase 一個 invocation 的平均 integer instructions , 取自 SQ_INSTS_VALU*64(wavefront size) ;
2. $R_{phase_i_avg}$: 各 phase 一個 invocation 的平均 shared mem/global mem 的 instructions , shared mem : SQ_INSTS_LDS*64*4 , global mem: SQ_INSTS_VMEM_RD/WR *64*4bytes ;
3. $t_{phase_i_avg}$: 各 phase 一個 invocation 的平均執行時間 , 單位為 ns ;

1. Blocking Factor:

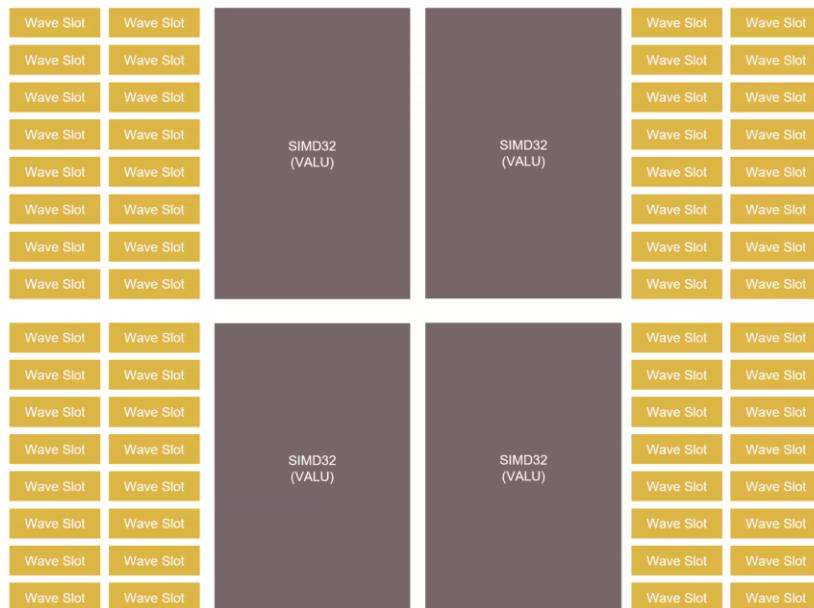


Testcase 為 c20.1 , Block_Factor 與 threads 數設定如下：

```
if [[ "$EXPERIMENT" == "hw3-2_smem-amd" ]]; then
    BLOCK_FACTORS=(16 32 48 64 80 90)
    THREAD_PER_BLOCK_X=(16 32 16 64 16 45)
    THREAD_PER_BLOCK_Y=(16 32 16 16 40 15)
```

由於這張 AMD GPU 的 shared memory per block 較大(65536 bytes)，所以 BF 可以設的比 GTX 1080 還大，可以看到，BF 超過 64 整體 GOPS 就開始下降了，可以看到即便 BF=80 跟 90 拉高了 ILP，也很難像 GTX 1080 一樣提升 performance，且相較於 NV 的 warp size 為 32，AMD 的 warp size(Wavefront)為 64，而 MI210 這張卡為 CDNA2 架構，每個 CU 包含 4 組 16-wide 的 SIMD(而 NV 是 32 wide)，因此 VALU instructions 實際會分成 4 個 SIMD slice，以 4 個 cycle 完成一個 waveform 的 64 threads(下圖為 RDNA 架構做為參考)，也就是說，只要確保每個 SIMD16 slice (即 waveform 的 0~15、16~31、32~47、48~63 這四組 lanes) 各自內部的 threads 都沒有 bank conflicts，那麼整個 waveform 就不會有 bank conflicts。

此為 RDNA 架構，SIMD 為 32-wide:



source: <https://gpuopen.com/learn/occupancy-explained/>

- For the AMD GPU, perform **one optimization** that was not done on the NVIDIA GPU, **or modify one NVIDIA-specific optimization** to suit the AMD GPU. Explain and show why this optimization is effective or ineffective, and clarify why it is specific to the AMD GPU.

基於 CDNA2 架構 16-wide SIMD 的架構，我們可以在現行的 access pattern 下，將 BF 及 blockDim.x 皆設為 16 的倍數，便可確保每個 SIMD instructions 的 shared mem access 不會跨 row，也就不會有 bank conflicts 發生：

BF=32, threads=32*32:

1	NSTS_VMEM_RD	"SQ_INSTS_VALU"	"SQ_INSTS_SALU"	"VFetchInsts"	"VWriteInsts"	"SQ_LDS_BANK_CONFLICT"
2	16,	1232,	32,	0,	0,	0,
3	9984,	1208096,	2795632,	0,	0,	0,
4	1168128,	30768704,	3145056,	0,	0,	0,
5	16,	1232,	32,	0,	0,	0,
6	9984,	1208096,	2795632,	0,	0,	0,

BF=48, threads=16*16:

```
experiments > amd > blocking_factor > rocprof > hw3-2_smem-amd > c20.1 > tmp_48.csv > data
1 , "SQ_INSTS_VMEM_RD", "SQ_INSTS_VALU", "SQ_INSTS_SALU", "VFetchInsts", "VWriteInsts", "SQ_LDS_BANK_CONFLICT",
2 , 36, 3572, 16, 0, 0, 0,
3 , 14976, 926828, 183056, 0, 0, 0,
4 , 1168128, 34524672, 524184, 0, 0, 0,
5 , 36, 3572, 16, 0, 0, 0,
6 , 14976, 926828, 183056, 0, 0, 0,
```

BF=80, threads=16*40:

```
experiments > amd > blocking_factor > rocprof > hw3-2_smem-amd > c20.1 > tmp_80.csv > data
1 "SQ_INSTS_VMEM_WR", "SQ_INSTS_VMEM_RD", "SQ_INSTS_VALU", "SQ_INSTS_SALU", "VFetchInsts", "VWriteInsts", "SQ_LDS_BANK_CONFLICT",
2 100, 100, 17260, 40, 0, 0, 0,
3 12400, 24800, 2209710, 172480, 0, 0, 0,
4 384400, 1153200, 52393720, 507220, 0, 0, 0,
5 100, 100, 17260, 40, 0, 0, 0,
6 12400, 24800, 2209710, 172480, 0, 0, 0,
```

BF=90, threads=45*15:

```
experiments > amd > blocking_factor > rocprof > hw3-2_smem-amd > c20.1 > tmp_90.csv > data
1 LWR", "SQ_INSTS_VMEM_RD", "SQ_INSTS_VALU", "SQ_INSTS_SALU", "VFetchInsts", "VWriteInsts", "SQ_LDS_BANK_CONFLICT",
2 132, 132, 24299, 55, 0, 0, 41328,
3 14520, 29040, 2876808, 152504, 0, 0, 3381840,
4 19300, 1197900, 69944050, 439901, 0, 0, 8639400,
5 132, 132, 24299, 55, 0, 0, 41328,
6 14520, 29040, 2876808, 152504, 0, 0, 3381840,
```

BF=32, threads=8*32:

```
experiments > amd > blocking_factor > rocprof > hw3-2_smem-amd > c20.1 > tmp_32.csv > data
1 STS_VMEM_WR", "SQ_INSTS_VMEM_RD", "SQ_INSTS_VALU", "SQ_INSTS_SALU", "VFetchInsts", "VWriteInsts", "SQ_LDS_BANK_CONFLICT",
2 16, 16, 1088, 8, 0, 0, 8736,
3 4992, 9984, 562856, 697660, 0, 0, 2006784,
4 389376, 1168128, 25212896, 883608, 0, 0, 10902528,
5 16, 16, 1088, 8, 0, 0, 8736,
6 4992, 9984, 562856, 697660, 0, 0, 2006784,
7 389376, 1168128, 25212896, 883608, 0, 0, 10902528,
```

當 blockDim.x 為 8 時因為不滿 16，所以會跨 2 個 rows，導致 2-way bank conflict。

另外，有發現在 AMD 上 thread coarsening 的效益不如在 NVIDIA 上明顯，主要可能的原因因為 AMD CDNA2 架構 CU 上一個 wave 執行一個 VALU instructions 需要 4 個 SIMD cycles 去執行，所以降低 occupancy 可能會導致 CU 上沒有足夠的 waves 給 warp scheduler 做切換來掩蓋 latency，透過 SQ_waves 能算出 occupancy:

BF=64, threads=64*16:

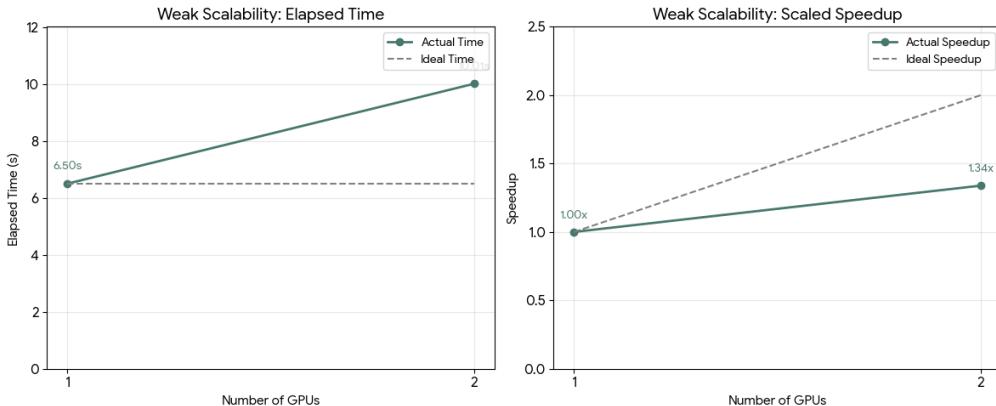
```
experiments > amd > blocking_factor > rocprof > hw3-2_smem-amd > c20.1 > tmp_64.csv > data
1 "arch_vgpr", "accum_vgpr", "sgpr", "wave_size", "sig", "obj" , "Wavefronts", "SQ_WAVES",
2 20, 28, 16, 64, "0x0", "0x153695130b40", 16, 16,
3 28, 28, 16, 64, "0x0", "0x153695130b80", 99856, 99856,
4 32, 32, 16, 64, "0x0", "0x153695130bc0", 99856, 99856,
5 20, 20, 16, 64, "0x0", "0x153695130b40", 16, 16,
6 28, 28, 16, 64, "0x0", "0x153695130b80", 99856, 99856,
```

BF=80, threads=16*40:

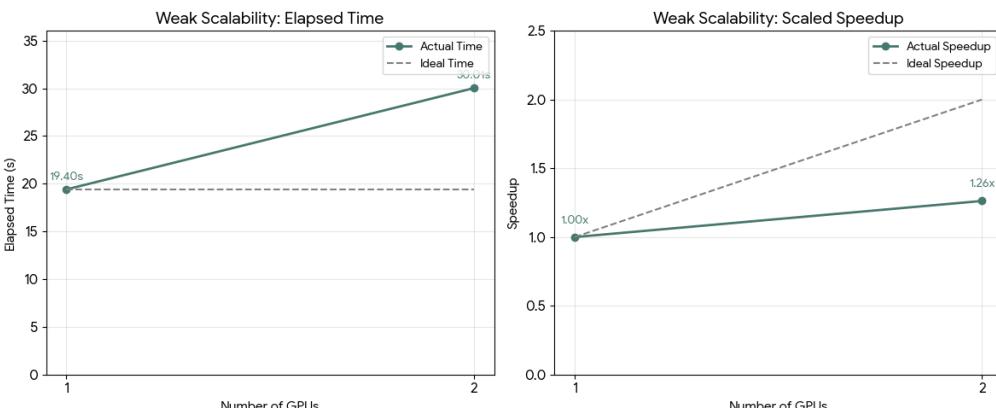
```
experiments > amd > blocking_factor > rocprof > hw3-2_smem-amd > c20.1 > tmp_80.csv > data
1 "1ds", "scr", "arch_vgpr", "accum_vgpr", "sgpr", "wave_size", "sig", "obj" , "Wavefronts", "SQ_WAVES",
2 10, 25600, 0, 36, 36, 16, 64, "0x0", "0x150d01020b40", 10, 10,
3 10, 51200, 0, 36, 36, 16, 64, "0x0", "0x150d01020b80", 39690, 39690,
4 10, 51200, 0, 56, 56, 16, 64, "0x0", "0x150d01020bc0", 39690, 39690,
5 10, 25600, 0, 36, 36, 16, 64, "0x0", "0x150d01020b40", 10, 10,
6 10, 51200, 0, 36, 36, 16, 64, "0x0", "0x150d01020b80", 39690, 39690,
```

3. According to Section 3-d, also conduct a **weak-scaling** experiment on the AMD GPU, and compare as well as discuss the results between the NVIDIA and AMD GPUs.

p11k1 v.s. p14k1:



p20k1 v.s. p25k1:



透過 weak scalability 的實驗可看出，AMD 跟 Nvidia 在 multi-gpu 時都因 communication 的成本難以達到 ideal speedup，甚至 speedup 比 Nvidia 還稍差一些，可以看到 GPU 間為透過 PCIE 溝通，但 hops 數為 2，應該是類似像 Nvidia PHB 有經過 CPU，所以 overhead 大：

```
pp25s023@apollo-login:~/hw3$ srun -p amd -N1 -n1 --gres=gpu:2 rcm-smi --showtopo

=====
===== ROCm System Management Interface =====
===== Weight between two GPUs =====
GPU0   GPU1
GPU0   0      40
GPU1   40      0

===== Hops between two GPUs =====
GPU0   GPU1
GPU0   0      2
GPU1   2      0

===== Link Type between two GPUs =====
GPU0   GPU1
GPU0   0      PCIE
GPU1   PCIE    0

===== Numa Nodes =====
GPU[0]      : (Topology) Numa Node: 1
GPU[0]      : (Topology) Numa Affinity: 1
GPU[1]      : (Topology) Numa Node: 1
GPU[1]      : (Topology) Numa Affinity: 1
=====
===== End of ROCm SMI Log =====
```

Experience & conclusion

a. What have you learned from this homework?

透過這次作業實作了多種 GPU programming 的優化技巧，深刻體會到 memory access 對效能的影響至關重要，其中像是 shared memory access pattern、global memory coalescing、BLOCK_FACTOR 的選擇等，都會極大影響 cuda 程式的效能，而 profiler 讓我在優化的過程中可以清楚看出程式哪邊還需要優化，以及讓我注意到了一些做 lab 時沒注意到的細節，此外，我也發現在使用 shared memory 的情況下，occupancy 並不是越大越好，降低 threads 數並適度地提高每個 thread 的工作量也可以提升整體效率；Multi-GPU 的優化則是比預期困難，GPU 間資料交換極大影響了整體的效能，很難達到 ideal 加速。